

# Heterogeneous Computing with SYCL

Kris Rowe

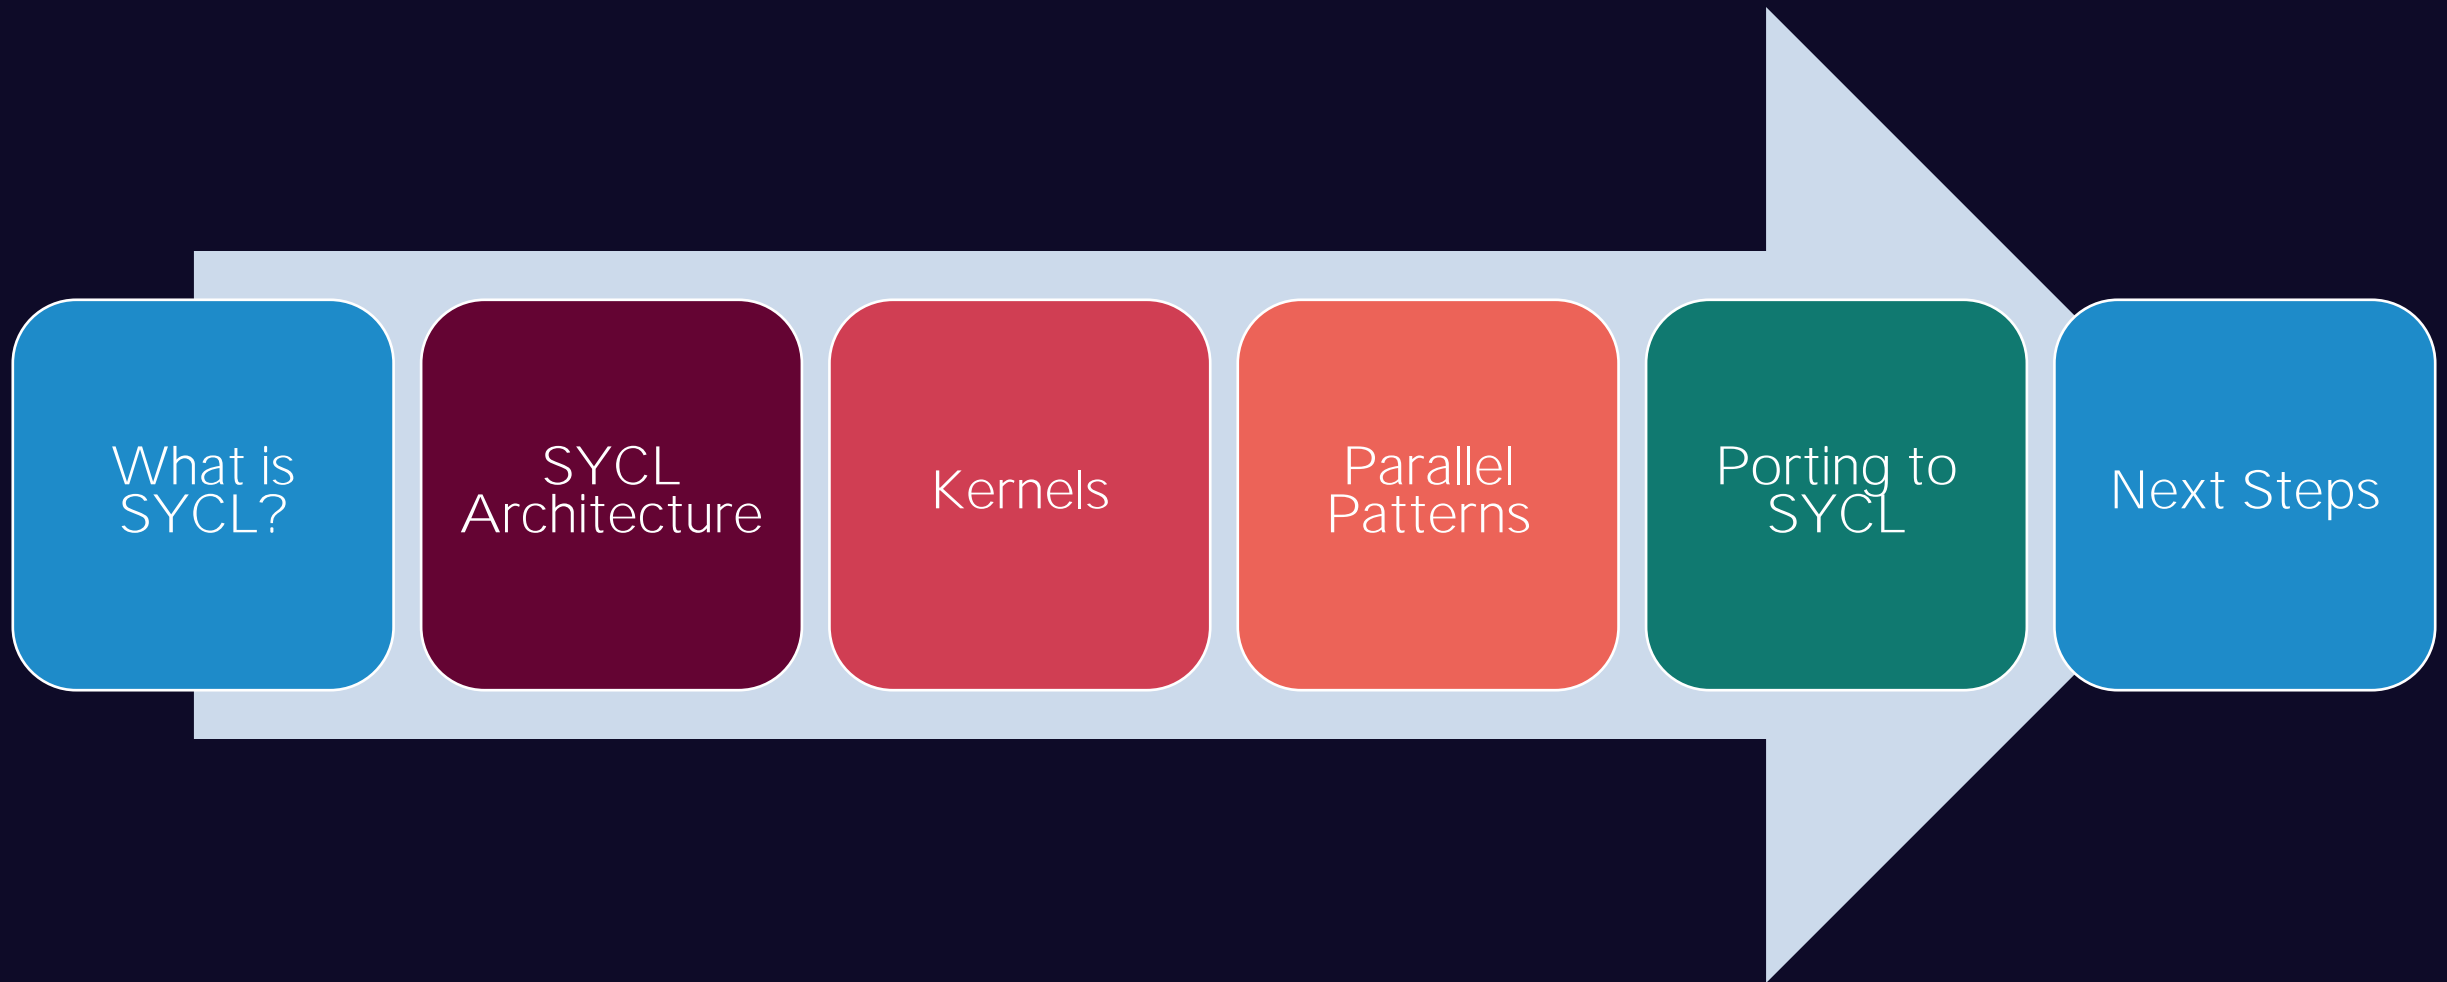
*Argonne Leadership Computing Facility*

Programming GPUs | 2022 Compute Ontario Summer School | July 8, 2022

# Course Material on GitHub

- <https://github.com/kris-rowe/coss-2022-sycl-tutorial>
- Code is not production grade, is meant for learning
- Examples are flat-coded for readability
- Exercises are somewhat refactored to hide boilerplate code
- Ask for help in GitHub discussions incase someone else has a similar question

# The Big Picture



# Approach

## USM instead of buffers/accessors

- Most HPC applications will use MPI
- Memory management is like CUDA
- Generally, provides better performance
- Explicit expression of data-dependencies
- Easily use libraries

## Lambdas instead of functors

- Source code is easier to understand
- Simpler to port existing code
- Kernels are defined where they are used

# What is SYCL?

# What is SYCL?



- An open, royalty-free, cross-platform standard for parallel programming on heterogeneous systems
- Developed by the Khronos Group
- A **single-source** programming model for modern standard C++ (C++17)
- Builds on concepts, portability, and efficiency of OpenCL
- *However, SYCL >> OpenCL + C++17*
  - Uses concepts that make modern C++ great
  - Higher-level interface is much more approachable than the OpenCL API
- Used by industry, government, and academia internationally
- Applications on systems ranging from laptops to world's largest supercomputers
- Runs on **different types of hardware**
  - CPUs
  - GPUs
  - FPGAs
- Runs on **different vendor's hardware**
  - NVIDIA
  - Intel
  - AMD

# Example SYCL Program

```
1  #include <CL/sycl.hpp>
2  #include <iostream>
3  #include <vector>
4
5  int main() {
6      const size_t vector_length = 2000;
7      std::vector<float> a_host(vector_length);
8      std::vector<float> b_host(vector_length, 1.0);
9      std::vector<float> c_host(vector_length, 1.0);
10     const float s = 1.0;
11
12     // Create a sycl::queue using the default device selector
13     sycl::device sycl_device{sycl::default_selector()};
14     sycl::context sycl_context{sycl_device};
15     sycl::queue sycl_queue{sycl_context, sycl_device};
16
17     // Allocate vectors on the device
18     float* a =
19         sycl::malloc_device<float>(vector_length, sycl_device, sycl_context);
20     float* b =
21         sycl::malloc_device<float>(vector_length, sycl_device, sycl_context);
22     float* c =
23         sycl::malloc_device<float>(vector_length, sycl_device, sycl_context);
24
25     // Copy from the host to the device; synchronize.
26     sycl_queue.copy(b_host.data(), b, b_host.size());
27     sycl_queue.copy(c_host.data(), c, c_host.size());
28     sycl_queue.wait();
29
30     // Submit work to the queue using a kernel defined via Lambdas; synchronize.
31     sycl_queue.parallel_for({vector_length},
32                             [=](sycl::id<1> i) { a[i] = b[i] + s * c[i]; });
33     sycl_queue.wait();
34
35     // Copy from the device to the host; synchronize.
36     sycl_queue.copy(a, a_host.data(), a_host.size());
37     sycl_queue.wait();
38
39     // Verify the results.
40     for (const auto& a_i : a_host) {
41         // Don't check for equality of floating-point values in production code!
42         if (2.0 != a_i) {
43             std::cout << "Verification failed!\n";
44             return EXIT_FAILURE;
45         }
46     }
47
48     std::cout << "Success!\n";
49
50     // Free device memory
51     sycl::free(a, sycl_context);
52     sycl::free(b, sycl_context);
53     sycl::free(c, sycl_context);
54     return EXIT_SUCCESS;
55 }
```

# Compiling, Linking, & Running Code

- Since SYCL programs use single-source C++ code building a program is straightforward
- To use LLVM clang and friends, pass the `-fsycl` flag
- Device functions should be in the same translation unit as the kernels which use them
- Like CUDA compilation, a fat-binary is produced
  - Contains binaries and/or low-level (SPIR-V, ptx) code needed to run on one or more target platforms
  - Enables cross-compilation, good for HPC systems
  - Can prescribe specific targets using triples like `-fsycl-targets=nvptx64-nvidia-cuda`
  - Ahead-of-time compilation is available for some backends via linker flags (e.g., `-Xsycl`)
- Consult compiler documentation for a full set of options
  - <https://intel.github.io/llvm-docs/UsersManual.html>



# SYCL Architecture

# Backends

- A SYCL Backend implements the SYCL programming model
- Typically, an implementation uses low-level, highly-optimized, vendor-specific drivers
- Intel oneAPI DPC++ Compiler comes with several backends
  - **OpenCL** – for Intel CPUs and GPUs using the Intel OpenCL Driver
  - **Level Zero** – for Intel GPUs using the Intel oneAPI Level Zero Driver
  - **CUDA** – for NVIDIA GPUs using the NVIDIA CUDA Driver
  - **HIP** – for AMD GPUs using the AMD ROCm Driver
- Direct interaction with SYCL backend should not be necessary for most applications
- Interoperability functions provide an interface to backend objects
  - Use case: calling a vendor-specific library in a SYCL application

# Platforms & Devices

- A SYCL device represents hardware which can execute SYCL kernels
- Three main types
  - CPU
  - GPU
  - Accelerator (FPGA)
- A SYCL platform is a collection of devices managed by a single backend
- The same hardware can appear as different devices in different platforms
- E.g., a single Intel GPU could appear in both the Level-Zero and OpenCL platforms

```
Platform 0: Intel(R) OpenCL
Device 0: Intel(R) Xeon(R) CPU E3-1585 v5 @ 3.50GHz
Type: CPU
Memory: 62 GB
Max Work Group Size: 8192

Platform 1: Intel(R) Level-Zero
Device 0: Intel(R) Iris(R) Pro Graphics P580 [0x193a]
Type: GPU
Memory: 49 GB
Max Work Group Size: 256

Platform 2: SYCL host platform
Device 0: SYCL host device
Type: host
Memory: 62 GB
Max Work Group Size: 1
```

A single CPU system with integrated graphics

# Information Descriptors

SYCL runtime classes provide the function `get_info`

The return value depends on the template parameter—called an information descriptor

```
for (auto& p : platforms) {
    std::string platform_name = p.get_info<sycl::info::platform::name>();
    std::cout << "Platform " << platform_id << ": " << platform_name << "\n";

    auto devices = p.get_devices();
    int device_id = 0;
    for (auto& d : devices) {
        std::string device_name = d.get_info<sycl::info::device::name>();
        std::cout << "Device " << device_id << ": " << device_name << "\n";

        printDeviceType(d);

        uint64_t memory = d.get_info<sycl::info::device::global_mem_size>();
        std::cout << "Memory: " << (memory / gigabyte) << " GB\n";

        uint64_t max_wg_size =
            d.get_info<sycl::info::device::max_work_group_size>();
        std::cout << "Max Work Group Size: " << max_wg_size << "\n";
    }
}
```

# Contexts

- A SYCL context is
  - A collection of one or more SYCL devices
  - Associated memory allocations
  - Any work-queues on each device
- This defines the fundamental “world” in which SYCL programs operate
- Convenient shortcuts avoid the need to deal with platforms and contexts explicitly

# Devices Selectors

- A device selector is a function which returns an integral score for a given SYCL device
- When passed to a constructor, the SYCL runtime scores all available SYCL devices and uses the highest-scoring device
- Several built-in device selectors are defined by the SYCL standard
  - `default_selector`
  - `cpu_selector`
  - `gpu_selector`

```
// Create a sycl::queue using the default device selector  
sycl::device sycl_device{sycl::default_selector()};  
sycl::context sycl_context{sycl_device};  
sycl::queue sycl_queue{sycl_context, sycl_device};
```

Exercise #2 involves implementing a custom device selector!

# Memory Management

- The SYCL Unified Shared Memory model uses a single virtual address space to ensure pointers in host and device code have consistent values.
- C-style malloc functions provide three flavors of memory allocation
  - **Device allocations**
    - ♦ Exist in on-device memory
    - ♦ Accessible on the device, but not the host
    - ♦ Data must be explicitly transferred between the host and device
  - **Host allocations**
    - ♦ Page-locked memory on the host
    - ♦ Accessible on both the host and device
    - ♦ Device can write to host memory directly in kernels
  - **Shared allocations**
    - ♦ Accessible on both the host and device
    - ♦ Automatically migrated by the runtime
- USM allocations must be explicitly freed by calling **sycl::free**

# Queues

- SYCL queues are work-queues on a specific device to which work can be submitted
- By default, SYCL queues are out-of-order
- No need to manage multiple in-order queues
  - Runtime handles the hard-work for you!
- Need to ensure data transfer operations are finished before they are needed by a kernel
- If necessary, can force queue to be in-order
  - Useful for porting CUDA code
- Most common functions that enqueue work
  - `parallel_for`
  - `memcpy`



# Host Synchronization & Events

- Constructor and malloc calls are synchronous
- Functions that submit work to a device are asynchronous, but return SYCL event objects
  - E.g., `memcpy` and `parallel_for`
- To synchronize with host
  - Call `queue.wait()` or `wait_and_throw()`
    - ♦ This waits for all enqueued work to finish—could block other queues with dependencies
  - Wait on a specific event (or set of events)
- Events can also be passed to work enqueueing functions as dependencies
  - Enables the programmer to directly express data dependencies

# Exceptions

SYCL defines an exception class that extends C++ standard exceptions

Exception handling is done with the usual **try-catch** pattern

```
1  #include <CL/sycl.hpp>
2  #include <iostream>
3
4  int main() {
5      try {
6          // Create a host device
7          sycl::device sycl_device{sycl::default_selector()};
8          sycl::context sycl_context{sycl_device};
9          sycl::queue sycl_queue{sycl_context, sycl_device};
10
11         int i{};
12         // Mock device_malloc failing and returning nullptr
13         int* j = nullptr;
14         sycl_queue.copy<int>(&i, j, 1);
15
16     } catch (const sycl::exception& e) {
17         std::cerr << "Synchronous exception\n";
18         std::cerr << e.what() << std::endl;
19     }
20     std::cout << "Exception handled.\n";
21     return EXIT_SUCCESS;
22 }
23
```

# Kernels

# Data Parallelism




A kernel is a function defined for a ranges of values in an index space

Instances of the function body are executed in parallel for each point in this space

While each instance, or work-item, executes the same code, the data operated on can vary

The mapping of work-items to hardware processing elements is implementation specific

# Thread Hierarchy



**Work-item**

# Defining Kernels Using Lambdas

- Kernels can be defined using C++ lambdas
- Basic kernels take `sycl::id` or `sycl::item` argument
- ND-range kernels take a `sycl::nd_item` argument
- Key difference: ND-range kernels can use shared local memory and synchronize work-items within the same work-group

```
[=](sycl::id<1> i) {  
    a[i] = b[i] + s * c[i];  
}
```

Basic Kernel

```
[=](sycl::nd_item<2> work_item) {  
    //...  
}
```

ND-Range Kernel

# Launching Kernels & Data Dependencies

```
sycl_queue.parallel_for({vector_length},  
    [=](sycl::id<1> i) {  
        a[i] = b[i] + s * c[i];  
    });
```

Basic Kernel

```
sycl::range<2> local_range(block_size, block_size);  
sycl::range<2> global_range(N, M);  
sycl::nd_range<2> kernel_range(global_range, local_range);  
  
sycl::event gemm_kernel = sycl_queue.parallel_for(  
    kernel_range, {copy_a, copy_b},  
    [=](sycl::nd_item<2> work_item) {  
        //...  
    });
```

ND-range Kernel

# Device Functions

- Since SYCL uses single-source C++ code, functions may be compiled for the host, the device, or both
- Extra restrictions are placed on device functions compared to regular C++ code
  - Data types must be trivially copyable
  - No memory allocation
  - No virtual functions
  - No RTTI
  - No variadic templates
  - No exception handling
- See the SYCL spec for a complete list:  
<https://www.khronos.org/registry/SYCL/specs/sycl-2020/html/sycl-2020.html#sec:language.restrictions.kernels>

```
1  void f(handler& cgh) {  
2      // Function "f" is not compiled for device  
3  
4      cgh.single_task( [= ] {  
5          // This code is compiled for device  
6          g(); // This line forces "g" to be compiled for device  
7      });  
8  }  
9  
10 void g() {  
11     // Called from kernel, so "g" is compiled for device  
12 }  
13  
14 void h() {  
15     // Not called from a device function, so not compiled for device  
16 }
```



# Parallel Patterns (Examples Demo)

# Porting to SYCL

# Porting an Existing C++ Application

- Modular, testable code is easier to port
  - Worth spending time improving test coverage, refactoring before porting
- Profile performance under realistic workloads relevant to your target science problem
- Identify parts of application that are most likely to benefit from using GPUs
  - E.g., Use Intel Advisor Offload Modelling  
<https://www.intel.com/content/www/us/en/develop/documentation/advisor-user-guide/top/model-offloading-to-a-gpu.html>
- Extract proxy/mini-apps which exercise a critical part of your code in isolation
- Port mini-apps, then incorporate ported code back into main application

# Porting an Existing CUDA Application

- Using USM approach manual code porting should be straightforward
- Start with in-order queues and ensure program correctness before relaxing to out-of-order
- Port code in stages using interoperability interfaces
- Automated porting tools exist—like Intel's DPC++ Compatibility Tool  
<https://www.intel.com/content/www/us/en/developer/tools/oneapi/dpc-compatibility-tool.html>
- Still requires intervention by the developer

# Next Steps

# Exercises

Several exercises are outlined on GitHub

Stubs or significant source code is provided for most of the exercises

Focus is on using SYCL concepts

<https://github.com/kris-rowe/coss-2022-sycl-tutorial/tree/main/exercises#exercises>

## Exercises

Complete the exercises you are most comfortable with first. Feeling up for a challenge? Try tackling some of the more difficult tasks. Need help or want to know if you are on the right track? Ask a question in the [Q&A discussions category](#).

The [SYCL Reference Guide](#) (cheat sheet) provides a concise summary of commonly used SYCL functions and is a helpful resource when first learning SYCL programming.

### 1. More Device Info

Extend the `device_info` example to provide more information about the available hardware. See the SYCL 2020 specification for a complete list of [device information descriptors](#).

Some types of hardware support extensions that are not available otherwise. Extensions include support for certain floating-point types, atomic operations, or memory allocation types. These extensions can be queried through the function

```
class device {  
public:  
    bool has(aspect asp) const
```

# Portable Libraries

oneMKL Interfaces

<https://github.com/oneapi-src/oneMKL>

Checkout other libraries defined as part of the oneAPI spec

<https://www.oneapi.io/spec/>

## oneAPI Math Kernel Library (oneMKL) Interfaces



oneMKL Interfaces is an open-source implementation of the oneMKL Data Parallel C++ (DPC++) interface according to the [oneMKL specification](#). It works with multiple devices (backends) using device-specific libraries underneath.

oneMKL is part of [oneAPI](#).

User Application	oneMKL Layer	Third-Party Library	Hardware Backend
oneMKL interface	oneMKL selector	<a href="#">Intel(R) oneAPI Math Kernel Library for x86 CPU</a>	x86 CPU
		<a href="#">Intel(R) oneAPI Math Kernel Library for Intel GPU</a>	Intel GPU
		<a href="#">NVIDIA cuBLAS for NVIDIA GPU</a>	NVIDIA GPU
		<a href="#">NVIDIA cuSOLVER for NVIDIA GPU</a>	NVIDIA GPU
		<a href="#">NVIDIA cuRAND for NVIDIA GPU</a>	NVIDIA GPU
		<a href="#">NETLIB LAPACK for x86 CPU</a>	x86 CPU
		<a href="#">AMD rocBLAS for AMD GPU</a>	AMD GPU
		<a href="#">AMD rocRAND for AMD GPU</a>	AMD GPU

# Resources

## Intel

- DPC++ Essentials Training
  - <https://www.intel.com/content/www/us/en/developer/tools/oneapi/training/dpc-essentials.html>
- GPU Optimization Guide
  - <https://www.intel.com/content/www/us/en/develop/documentation/oneapi-gpu-optimization-guide/top.html>
- oneAPI samples
  - <https://github.com/oneapi-src/oneAPI-samples>
- Developer Conferences
  - <https://www.oneapi.io/events/>
- DevMesh
  - <https://devmesh.intel.com/>

## ALCF

- Events
  - <https://www.alcf.anl.gov/events>
- Training
  - <https://www.alcf.anl.gov/support-center/training-assets>
- ATPESC
  - <https://extremecomputingtraining.anl.gov/>



# References

- SYCL 2020 Specification
  - <https://www.khronos.org/registry/SYCL/specs/sycl-2020/html/sycl-2020.html>
- SYCL Reference Guide (cheat sheet)
  - <https://www.khronos.org/files/sycl/sycl-2020-reference-guide.pdf>
- oneAPI DPC++ (Intel LLVM) Compiler Documentation
  - <https://intel.github.io/llvm-docs/>

This work was supported by Argonne Leadership Computing Facility, which is a DOE Office of Science User Facility supported under Contract DE-AC02-06CH11357.